

# Logistic Regression


Predicting probabilities with an “S” curve

Faculty of Mathematics and Computer Science, University of Bucharest  
and  
Sparktech Software

*Academic Year 2018/2019, 1<sup>st</sup> Semester*

# Classification vs. Regression

- Previously, we wanted to predict the price of a house, given its size.



# Classification vs. Regression


- Previously, we wanted to predict the price of a house, given its size.
- Now, we want to predict if a house costs more than \$450k or not.
  - Not only that, but we only have access to a set of samples with this information.



# Classification vs. Regression


- Previously, we wanted to predict the price of a house, given its size.
- Now, we want to predict if a house costs more than \$450k or not.
  - Not only that, but we only have access to a set of samples with this information.
  - Even better, we want a model which gives us the **likelihood** of a house costing more than \$450k

## Chances of Price > \$450k?



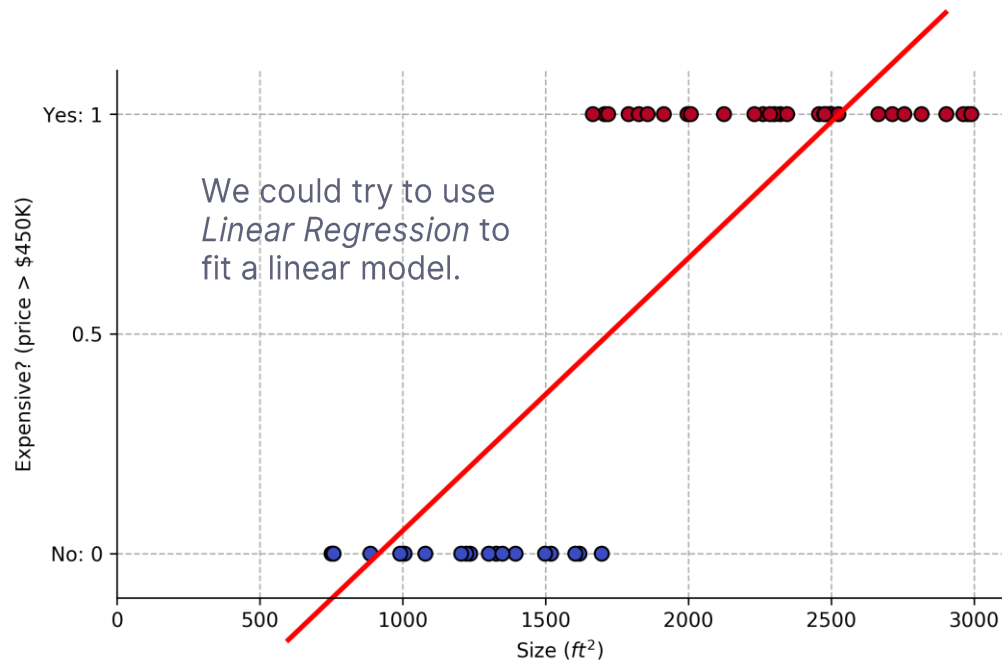
# Classification vs. Regression

- Set of *expensive / non-expensive* house observations:




# Classification vs. Regression

- Set of *expensive / non-expensive* house observations:




# Classification vs. Regression

- Set of *expensive / non-expensive* house observations.
  - Is there anything wrong with this approach?



# Classification vs. Regression

- Set of *expensive / non-expensive* house observations.
  - Is there anything wrong with this approach?

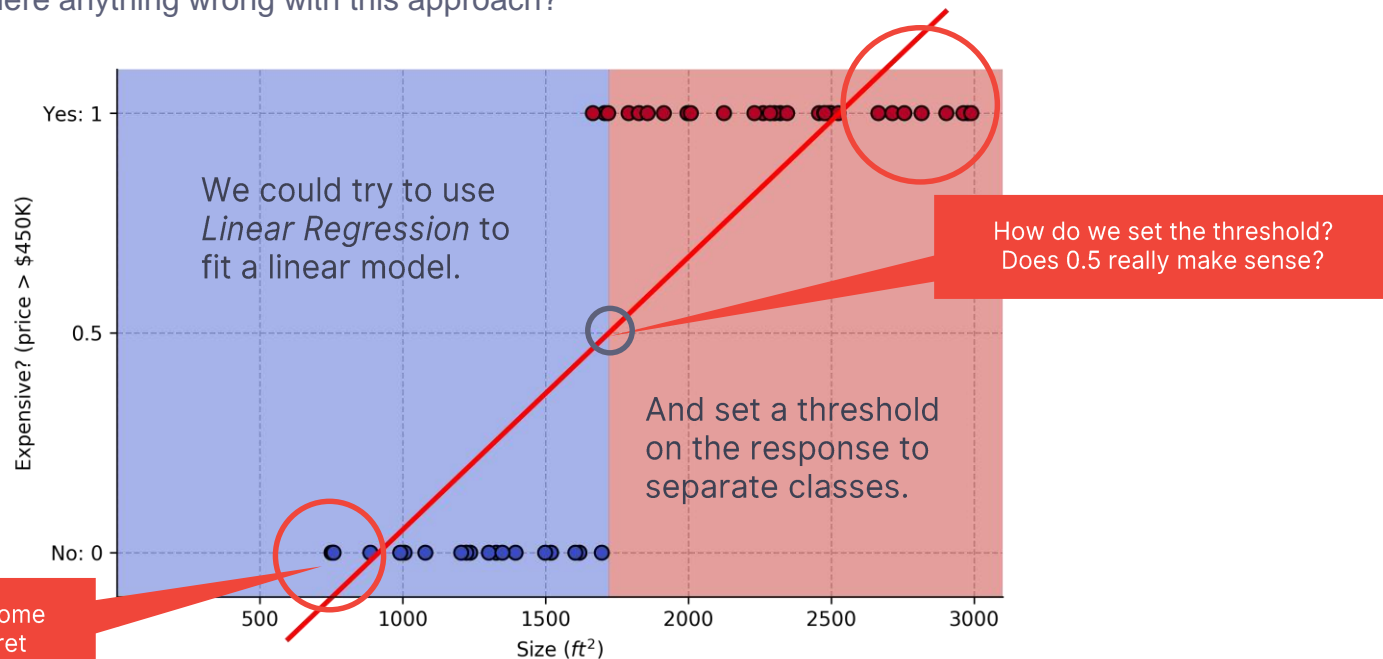


Some outputs are  $< 0$  and some are  $> 1$ , so we can't interpret them as probabilities




# Classification vs. Regression

- Set of *expensive / non-expensive* house observations.
  - Is there anything wrong with this approach?



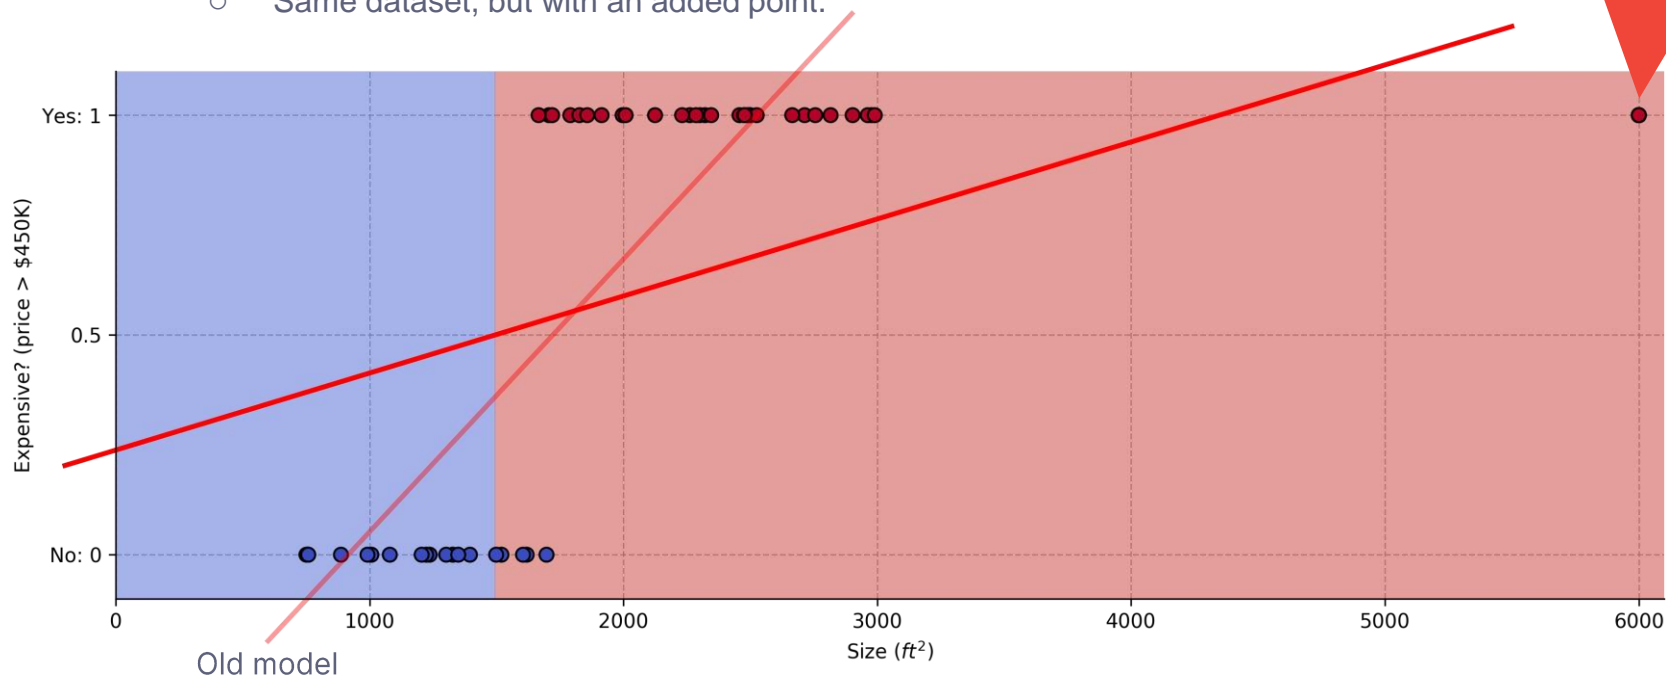
# Classification vs. Regression

- Set of *expensive / non-expensive* house observations.
  - Same dataset, but with an added point.




# Classification vs. Regression

- Set of *expensive / non-expensive* house observations.
  - Same dataset, but with an added point.




# Classification vs. Regression

- Set of *expensive / non-expensive* house observations.
  - Same dataset, but with an added point.




# Predicting probabilities

- Set of *expensive / non-expensive* house observations.
  - We want to learn a model  $h$ , which gives us the **probability** of a house being expensive, given its size.



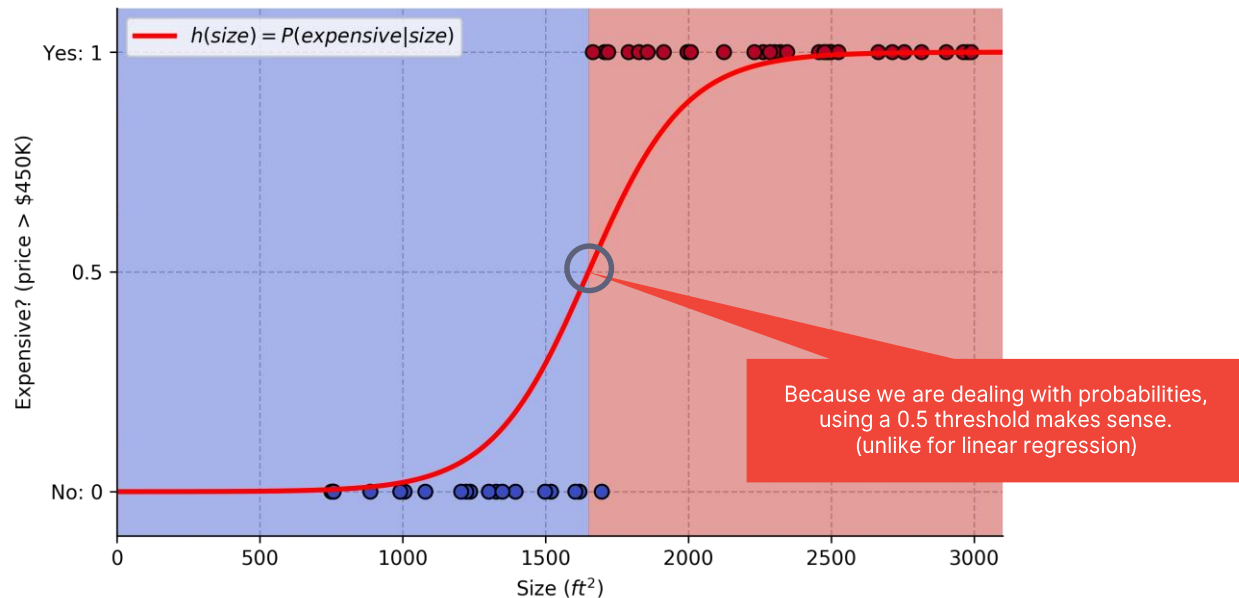
# Predicting probabilities

- Set of *expensive / non-expensive* house observations.
  - We want to learn a model  $h$ , which gives us the **probability** of a house being expensive, given its size.



# Predicting probabilities

- Set of *expensive / non-expensive* house observations.
  - We want to learn a model  $h$ , which gives us the **probability** of a house being expensive, given its size.



# Logistic Function


- Sigmoid function (“S-shaped” curve):
  - Family of functions which are bounded, differentiable, real and with a non-negative derivative at each point.
- Logistic function (special case of sigmoid):

$$\sigma(x) = \frac{L}{1 + e^{-k(x-x_0)}}$$

- Standard logistic function ( $L = 1, k = 1, x_0 = 0$ ):

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- $\sigma(-x) = 1 - \sigma(x) = \frac{1}{1+e^x}$
- $0 < \sigma(x) < 1 \Rightarrow$  Can be interpreted as **probability**





# Logistic Model

- The relationship is modeled with a *logistic model*:
  - $\vec{x} \in \mathbb{R}^n$  represents the independent variables
  - $y \in \{0, 1\}$  is the dependent variable

$$\hat{y} = \sigma(w_0 + w_1x_1 + \cdots + w_nx_n)$$

# Logistic Model

- The relationship is modeled with a *logistic model*:
  - $\vec{x} \in \mathbb{R}^n$  represents the independent variables
  - $y \in \{0, 1\}$  is the dependent variable

$$\hat{y} = \sigma(w_0 + w_1x_1 + \dots + w_nx_n) = \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}}$$

Same trick as  
before:  $x_0 = 1$

- In other words, we apply a logistic function on top of a linear model.

# Logistic Model

- The relationship is modeled with a *logistic model*:
  - $\vec{x} \in \mathbb{R}^n$  represents the independent variables
  - $y \in \{0, 1\}$  is the dependent variable

$$\hat{y} = \sigma(w_0 + w_1x_1 + \dots + w_nx_n) = \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}}$$

Same trick as  
before:  $x_0 = 1$

- In other words, we apply a logistic function on top of a linear model.

- Prediction can be interpreted as probability:

$$\hat{y} = P(y = 1 | \vec{x}, \vec{w})$$

# Loss function

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{0,1\}$$

# Loss function

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{0,1\}$$

- We could use *squared-error loss* like in linear regression.

$$\mathcal{L}_E = \sum_i (y^{(i)} - \hat{y}^{(i)})^2 = \sum_i \left( y^{(i)} - \frac{1}{1 + e^{-\langle \vec{w}, \vec{x}^{(i)} \rangle}} \right)^2$$

# Loss function

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{0,1\}$$

- We could use *squared-error loss* like in linear regression.

$$\mathcal{L}_E = \sum_i (y^{(i)} - \hat{y}^{(i)})^2 = \sum_i \left( y^{(i)} - \frac{1}{1 + e^{-\langle \vec{w}, \vec{x}^{(i)} \rangle}} \right)^2$$



- But this turns out to be a **non-convex function**

# Convex vs. Non-convex

- A function is **convex** if a *segment between* any two points on its graph lies *above* the graph.

# Convex vs. Non-convex

- A function is **convex** if a *segment between any two points on its graph* lies *above* the graph.





- $f: X \rightarrow Y$  is convex if  $f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2) \quad \forall x_1, x_2 \in X, \forall t \in [0, 1]$



# Convex vs. Non-convex

- A function is **convex** if a *segment between any two points on its graph* lies *above* the graph.



- $f: X \rightarrow Y$  is convex if  $f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2) \quad \forall x_1, x_2 \in X, \forall t \in [0, 1]$
- A convex function has *only one local minimum*  $\rightarrow$  We want error functions to be convex.

# Loss function

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{0,1\}$$

- We could use *squared-error loss* like in linear regression:

$$\mathcal{L}_E = \sum_i (y^{(i)} - \hat{y}^{(i)})^2 = \sum_i \left( y^{(i)} - \frac{1}{1 + e^{-\langle \vec{w}, \vec{x}^{(i)} \rangle}} \right)^2$$

- But this turns out to be a **non-convex function**  $\Rightarrow$  multiple local minima

# Loss function

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{0,1\}$$

- We could use *squared-error loss* like in linear regression:


$$\mathcal{L}_E = \sum_i (y^{(i)} - \hat{y}^{(i)})^2 = \sum_i \left( y^{(i)} - \frac{1}{1 + e^{-\langle \vec{w}, \vec{x}^{(i)} \rangle}} \right)^2$$

- But this turns out to be a **non-convex function**  $\Rightarrow$  multiple local minima
- Using the so-called **Cross-entropy** or **Log-loss** works better for logistic regression.

# Loss function

- Cross-entropy for a single example:


$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \begin{cases} -\log \hat{y}^{(i)} & \text{if } y^{(i)} = 1 \end{cases}$$



# Loss function

- Cross-entropy for a single example:

$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \begin{cases} -\log \hat{y}^{(i)} & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$




# Loss function

- Cross-entropy for a single example:

$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \begin{cases} -\log \hat{y}^{(i)} & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

- $\mathcal{L} \rightarrow 0$  as  $\hat{y}^{(i)} \rightarrow y^{(i)}$
- $\mathcal{L} \rightarrow \infty$  as  $\hat{y}^{(i)} \rightarrow 1 - y^{(i)}$

Loss grows exponentially if model is very confident in the wrong prediction.




# Loss function

- Cross-entropy for a single example:

$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \begin{cases} -\log \hat{y}^{(i)} & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

- $\mathcal{L} \rightarrow 0$  as  $\hat{y}^{(i)} \rightarrow y^{(i)}$
  - $\mathcal{L} \rightarrow \infty$  as  $\hat{y}^{(i)} \rightarrow 1 - y^{(i)}$
- Loss grows exponentially if model is very confident in the wrong prediction.



- Combining both branches and summing over all samples:


$$\mathcal{L}_E = - \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

# Loss function

- Cross-entropy for a single example:

$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \begin{cases} -\log \hat{y}^{(i)} & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

- $\mathcal{L} \rightarrow 0$  as  $\hat{y}^{(i)} \rightarrow y^{(i)}$
  - $\mathcal{L} \rightarrow \infty$  as  $\hat{y}^{(i)} \rightarrow 1 - y^{(i)}$
- Loss grows exponentially if model is very confident in the wrong prediction.



- Combining both branches and summing over all samples:

$$\mathcal{L}_E = - \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \lambda \|\vec{w}\|_2^2$$

Regularized Logistic Regression



# Optimization


- We need to find  $\vec{w}_* = \operatorname{argmin}_{\vec{w}} \mathcal{L}_E(\vec{w})$

# Optimization

- We need to find  $\vec{w}_* = \underset{\vec{w}}{\operatorname{argmin}} \mathcal{L}_E(\vec{w})$
- There is no *closed-form solution* for finding the minimum of the log-loss.
  - i.e. We cannot compute it directly through mathematical operations, like for linear regression
  - We need to use an *optimization method*.


# Optimization

- We need to find  $\vec{w}_* = \underset{\vec{w}}{\operatorname{argmin}} \mathcal{L}_E(\vec{w})$
- There is no *closed-form solution* for finding the minimum of the log-loss.
  - i.e. We cannot compute it directly through mathematical operations, like for linear regression
  - We need to use an *optimization method*.
- Gradient Descent
  - Gradually go down the slope of the error function



# Optimization

- We need to find  $\vec{w}_* = \underset{\vec{w}}{\operatorname{argmin}} \mathcal{L}_E(\vec{w})$
- There is no *closed-form solution* for finding the minimum of the log-loss.
  - i.e. We cannot compute it directly through mathematical operations, like for linear regression
  - We need to use an *optimization method*.
- Gradient Descent
  - Gradually go down the slope of the error function
- More complex methods:
  - Conjugate Gradient
  - BFGS
  - L-BFGS



## Optimization – $\mathcal{C}$ vs. $\lambda$

$$\operatorname{argmin}_{\vec{w}} \mathcal{L}_E(\vec{w}) =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ - \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \lambda \|\vec{w}\|_2^2 \right\}$$

# Optimization – $\mathcal{L}$ vs. $\lambda$

$$\operatorname{argmin}_{\vec{w}} \mathcal{L}_E(\vec{w}) =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ - \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \lambda \|\vec{w}\|_2^2 \right\} =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ - \frac{1}{\lambda} \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \|\vec{w}\|_2^2 \right\}$$

Dividing an expression by a constant does not change the *argmin*.

# Optimization – $\mathcal{C}$ vs. $\lambda$

$$\operatorname{argmin}_{\vec{w}} \mathcal{L}_E(\vec{w}) =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ - \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \lambda \|\vec{w}\|_2^2 \right\} =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ - \frac{1}{\lambda} \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \|\vec{w}\|_2^2 \right\} =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ -C \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \|\vec{w}\|_2^2 \right\}$$

Dividing an expression by a constant does not change the *argmin*.

$$1/\lambda \stackrel{\text{not}}{=} C$$

# Optimization – $C$ vs. $\lambda$

$$\operatorname{argmin}_{\vec{w}} \mathcal{L}_E(\vec{w}) =$$

$$\operatorname{argmin}_{\vec{w}} \left\{ - \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \lambda \|\vec{w}\|_2^2 \right\} =$$

Dividing an expression by a constant does not change the *argmin*.

$$\operatorname{argmin}_{\vec{w}} \left\{ - \frac{1}{\lambda} \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \|\vec{w}\|_2^2 \right\} =$$

$$1/\lambda \stackrel{\text{not}}{=} C$$

$$\operatorname{argmin}_{\vec{w}} \left\{ -C \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \|\vec{w}\|_2^2 \right\}$$

- More common formulation of regularization for classification problems, because  $C$  is easier to interpret (“*cost of making a training mistake*”).



# **Different Formulation for Loss Function**

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

$$P(y = 1 | \vec{x}, \vec{w}) = \sigma(\langle \vec{w}, \vec{x} \rangle) = \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}}$$

We can denote the two classes of binary classification with  $\pm 1$

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

$$P(y = 1 | \vec{x}, \vec{w}) = \sigma(\langle \vec{w}, \vec{x} \rangle) = \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}}$$

$$P(y = -1 | \vec{x}, \vec{w}) = 1 - \sigma(\langle \vec{w}, \vec{x} \rangle) = \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}}$$

We can denote the two classes of binary classification with  $\pm 1$

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

$$\left. \begin{aligned} P(y = 1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}} \\ P(y = -1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}} \end{aligned} \right\} P(y^{(i)} | \vec{x}^{(i)}, \vec{w}) = \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

$$\left. \begin{aligned} P(y = 1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}} \\ P(y = -1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}} \end{aligned} \right\} P(y^{(i)} | \vec{x}^{(i)}, \vec{w}) = \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

- Likelihood of data:

$$\prod_{i=1}^m P(y^{(i)}) = \prod_{i=1}^m \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

$$\left. \begin{aligned} P(y = 1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}} \\ P(y = -1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}} \end{aligned} \right\} P(y^{(i)} | \vec{x}^{(i)}, \vec{w}) = \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

- Log-Likelihood of data:  $\log \left( \prod_{i=1}^m \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}} \right) = - \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right)$

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

$$\left. \begin{aligned} P(y = 1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}} \\ P(y = -1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}} \end{aligned} \right\} P(y^{(i)} | \vec{x}^{(i)}, \vec{w}) = \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

- Log-Likelihood of data:  $\log \left( \prod_{i=1}^m \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}} \right) = - \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right)$
- Maximize the log-likelihood of data  $\Rightarrow$

$$\text{minimize } \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right)$$



# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

$$\left. \begin{aligned} P(y = 1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}} \\ P(y = -1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}} \end{aligned} \right\} P(y^{(i)} | \vec{x}^{(i)}, \vec{w}) = \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

- Log-Likelihood of data:  $\log \left( \prod_{i=1}^m \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}} \right) = - \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right)$
- Maximize the log-likelihood of data  $\Rightarrow$

$$\text{minimize } C \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right) + \|\vec{w}\|_2^2$$

With regularization

# Different Formulation

$$E = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})\}, \vec{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}$$

We can denote the two classes of binary classification with  $\pm 1$

$$\left. \begin{aligned} P(y = 1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}} \\ P(y = -1 | \vec{x}, \vec{w}) &= \frac{1}{1 + e^{\langle \vec{w}, \vec{x} \rangle}} \end{aligned} \right\} P(y^{(i)} | \vec{x}^{(i)}, \vec{w}) = \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}}$$

- Log-Likelihood of data:  $\log \left( \prod_{i=1}^m \frac{1}{1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle}} \right) = - \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right)$
- Maximize the log-likelihood of data  $\Rightarrow$

$$\text{minimize } C \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right) + \|\vec{w}\|_2^2$$

Formulation used by *Scikit-learn*.

With regularization

# Recap

- **Logistic Regression** uses a *logistic function* on top of a *linear model* to establish a relationship between a **binary** dependent variable and a number of *independent variables*.

$$\hat{y} = \sigma(\langle \vec{w}, \vec{x} \rangle) = \frac{1}{1 + e^{-\langle \vec{w}, \vec{x} \rangle}}$$

- Parameters are obtained by minimizing the **Cross-entropy loss**:

$$-C \sum_i [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] + \|\vec{w}\|_2^2$$

or by minimizing the negative **log-likelihood** of the data:

$$C \sum_{i=1}^m \log \left( 1 + e^{-y^{(i)} \langle \vec{w}, \vec{x}^{(i)} \rangle} \right) + \|\vec{w}\|_2^2$$

Considering labels to be  $\pm 1$

# **Multinomial Logistic Regression**

# What if we have multiple classes?

- We want to know *how likely it is* that a certain dog is of one of four breeds.



**Weight** = 25 kg  
**Height** = 46 cm  
**Fur length** = 0  
**Ear length** = 3 cm  
**Color** = Cream



German Sheppard



Labrador Retriever



Shar Pei



Collie

# What if we have multiple classes?

- We want to know *how likely it is* that a certain dog is of one of four breeds.



Weight = 25 kg  
Height = 46 cm  
Fur length = 0  
Ear length = 3 cm  
Color = Cream



# What if we have multiple classes?

- We want to know *how likely it is* that a certain dog is of one of four breeds.



Weight = 25 kg  
Height = 46 cm  
Fur length = 0  
Ear length = 3 cm  
Color = Cream



Predictions  
should sum  
up to 100%.

The predicted class  
is the one with the  
highest probability.

# What if we have multiple classes?

- If we have  $K$  classes, we predict a probability for each class:
  - The prediction  $\hat{y}^{(i)}$  becomes array of probabilities.
  - $\hat{y}_k^{(i)}$  is the probability of  $\vec{x}^{(i)}$  being class  $k$  (e.g.  $\hat{y}^{(i)} = [0.1 \quad 0.15 \quad 0.7 \quad 0.05]$ )



# What if we have multiple classes?

- If we have  $K$  classes, we predict a probability for each class:
  - The prediction  $\hat{y}^{(i)}$  becomes array of probabilities.
  - $\hat{y}_k^{(i)}$  is the probability of  $\vec{x}^{(i)}$  being class  $k$  (e.g.  $\hat{y}^{(i)} = [0.1 \quad 0.15 \quad 0.7 \quad 0.05]$ )

The predicted class is the one with the highest probability.

# What if we have multiple classes?

- If we have  $K$  classes, we predict a probability for each class:

- The prediction  $\hat{y}^{(i)}$  becomes array of probabilities.
- $\hat{y}_k^{(i)}$  is the probability of  $\vec{x}^{(i)}$  being class  $k$  (e.g.  $\hat{y}^{(i)} = [0.1 \quad 0.15 \quad 0.7 \quad 0.05]$ )

The predicted class is the one with the highest probability.

- Training samples:  $y_k^{(i)} = \begin{cases} 1 & \text{if } \vec{x}^{(i)} \text{ is class } k \\ 0 & \text{otherwise} \end{cases}$  (e.g.  $y^{(i)} = [0 \quad 0 \quad 1 \quad 0]$ )

# What if we have multiple classes?

- If we have  $K$  classes, we predict a probability for each class:

The predicted class is the one with the highest probability.

- The prediction  $\hat{y}^{(i)}$  becomes array of probabilities.
- $\hat{y}_k^{(i)}$  is the probability of  $\vec{x}^{(i)}$  being class  $k$  (e.g.  $\hat{y}^{(i)} = [0.1 \quad 0.15 \quad 0.7 \quad 0.05]$ )
- Training samples:  $y_k^{(i)} = \begin{cases} 1 & \text{if } \vec{x}^{(i)} \text{ is class } k \\ 0 & \text{otherwise} \end{cases}$  (e.g.  $y^{(i)} = [0 \quad 0 \quad 1 \quad 0]$ )

- **Multinomial Cross-Entropy:**

$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \sum_{k=1}^K y_k^{(i)} \log \hat{y}_k^{(i)}$$

# What if we have multiple classes?

- If we have  $K$  classes, we predict a probability for each class:

The predicted class is the one with the highest probability.

- The prediction  $\hat{y}^{(i)}$  becomes array of probabilities.
- $\hat{y}_k^{(i)}$  is the probability of  $\vec{x}^{(i)}$  being class  $k$  (e.g.  $\hat{y}^{(i)} = [0.1 \quad 0.15 \quad 0.7 \quad 0.05]$ )
- Training samples:  $y_k^{(i)} = \begin{cases} 1 & \text{if } \vec{x}^{(i)} \text{ is class } k \\ 0 & \text{otherwise} \end{cases}$  (e.g.  $y^{(i)} = [0 \quad 0 \quad 1 \quad 0]$ )

- **Multinomial Cross-Entropy:**

$$\mathcal{L}(y^{(i)}, \hat{y}^{(i)}) = \sum_{k=1}^K y_k^{(i)} \log \hat{y}_k^{(i)}$$

Only one term will be non-zero

# What if we have multiple classes?

- We have a separate weight vector  $\vec{w}_k$  for each class.
- Prediction is computed by applying **Softmax**:

$$\hat{y}_k = \frac{e^{\langle \vec{w}_k, \vec{x} \rangle}}{\sum_{j=1}^K e^{\langle \vec{w}_j, \vec{x} \rangle}}$$

*Softmax* is a generalization of the logistic function for n-dimensional inputs.

# What if we have multiple classes?

- We have a separate weight vector  $\vec{w}_k$  for each class.
- Prediction is computed by applying **Softmax**:

$$\hat{y}_k = \frac{e^{\langle \vec{w}_k, \vec{x} \rangle}}{\sum_{j=1}^K e^{\langle \vec{w}_j, \vec{x} \rangle}}$$

*Softmax* is a generalization of the logistic function for n-dimensional inputs.

- Even though we have a separate  $\vec{w}_j$  for each class, there is only one loss which is minimized.
  - The weight vectors are obtained simultaneously.

## Other strategies for multiclass classification

- Another method of having multiple classes is to fit multiple binary classifiers independently and combine their predictions.

# Other strategies for multiclass classification

- Another method of having multiple classes is to fit multiple binary classifiers independently and combine their predictions.
- **One-versus-rest** (OVR), sometimes called (one-versus-all, OVA)
  - Train  $n$  classifiers, one for each class, where the negative examples are all the other classes.
  - At inference, run all classifiers and pick the class with the highest margin (most confident)



# Other strategies for multiclass classification

- Another method of having multiple classes is to fit multiple binary classifiers independently and combine their predictions.
- **One-versus-rest (OVR)**, sometimes called (one-versus-all, OVA)
  - Train  $n$  classifiers, one for each class, where the negative examples are all the other classes.
  - At inference, run all classifiers and pick the class with the highest margin (most confident)
- **One-versus-one (OVO)**
  - Train  $n(n - 1)/2$  classifiers, one for each pair of classes.
  - At inference, run all classifiers and pick the class which was selected by most of them

# Logistic Regression in Python

```
1  from sklearn.linear_model import LogisticRegression
2
3  clf = LogisticRegression(C = 10) # C is the inverse regularization strength 1/λ
4  clf.fit(X, y)
5
6  clf.predict([x]) # prediction for x
7  clf.predict_proba([x]) # predicted probability for x
8  clf.decision_function([x]) # value of the decision function before applying logistic:  $\langle \vec{w}, \vec{x} \rangle$ 
9
10  clf.coef_ #  $w_1, w_2 \dots w_n$ 
11  clf.intercept_ #  $w_0$ 
12
13  clf = LogisticRegression(multi_class = 'multinomial') # multi_class = {'multinomial', 'ovr'}
```

# Conclusions

- **Logistic regression** uses a *logistic function* on top of a *linear function* to establish a relationship between a **binary dependent variable** and a number of *independent variables*.
- **Logistic Regression** is a method for **classification**.
  - Name is due to historical reasons and the relation to *linear regression*.
- The prediction can be interpreted as **probability**.
- The parameters of the model are obtained by minimizing the **cross-entropy loss** or *log-loss*.
  - Another possibility is to maximize the **log-likelihood** of the data
  - Since there is no *closed-form solution*, a convex optimization method is used.
- *Multinomial cross-entropy* can be used to achieve multiclass classification.
  - Other multiclass strategies are *OVR* and *OVO*.

# History of Logistic Regression

- *“The regression analysis of binary sequences (with discussion).”*

*David Cox, 1958*

- *“Estimation of the Probability of an Event as a Function of Several Independent Variables.”*

*Strother H. Walker and David B. Duncan, 1967*

# Keywords

Logistic Regression

Sigmoid

Probability

Cross-entropy

Closed-form Solution

Log-likelihood

Convex Function

Optimization Method

Gradient Descent

Multinomial

Softmax

One-versus-rest (OVR)

One-versus-one (OVO)